

Instant Cup Noodles

Condition Detector

By: Bryan (P2319104)
Ei Zin (P2346438)

Problem statement

Why this project?

- Visual differences between cup noodle states can be subtle, and yet
- Manual inspection is:
 - Time-consuming
 - Inconsistent & Tedious
- A computer vision model can:
 - Learn visual patterns (cup noodles have a distinct shape)
 - Classify conditions consistently
 - Work in real time



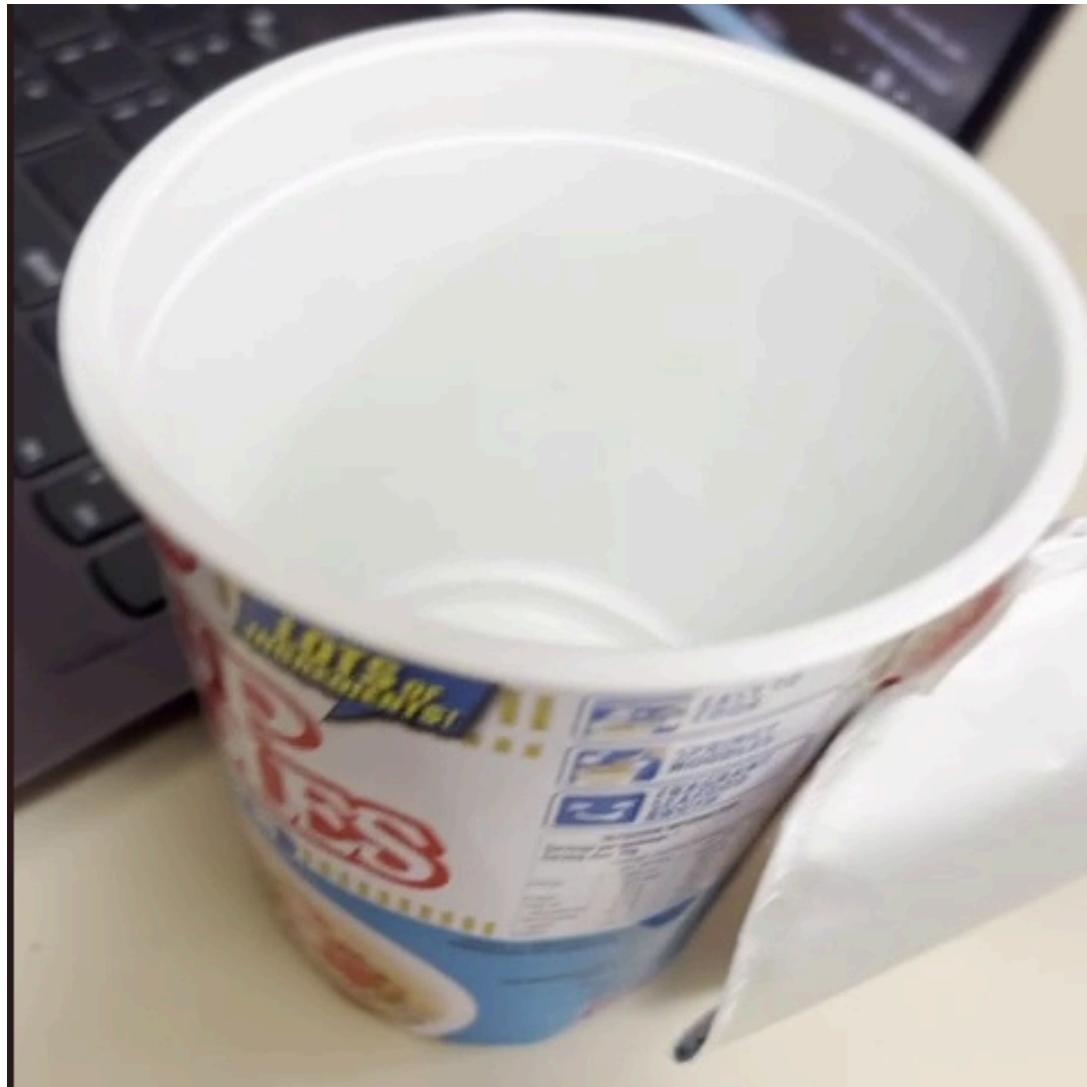
Project Objective

Objective

- Detect the condition of instant cup noodles
- The conditions of the cup noodles are:
 - Sealed
 - Opened
 - Finished
- Support image testing and live camera detection.
- Be able to detect the condition reliably across multiple brands and sizes of cup noodles.



Classes



Finished



Opened



Sealed

Data Collection

**How the dataset
was created**

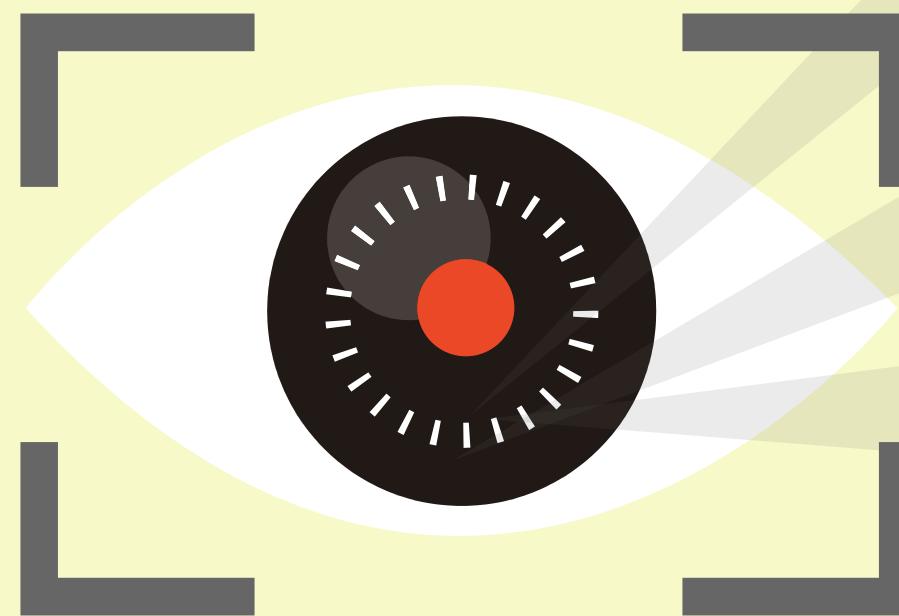


Image captured

Mobile phone camera & video

**Multiple
environments**

- Different lighting conditions
- Different backgrounds
- Different angles

Dataset

Dataset prepared manually
and labelled into 3 classes

Data Collection

```
cap = cv2.VideoCapture(video_path) # Opens video file

while True:
    ret, frame = cap.read() # read NEXT frame
    if not ret:
        break # no more frames in the video

    if frame_idx % step == 0:
        cv2.imwrite(f"{output_dir}/{filename}_{saved_idx}.jpg", frame)
        saved_idx += 1

    frame_idx += 1
```



.mov (video)
turned into .jpg (images)



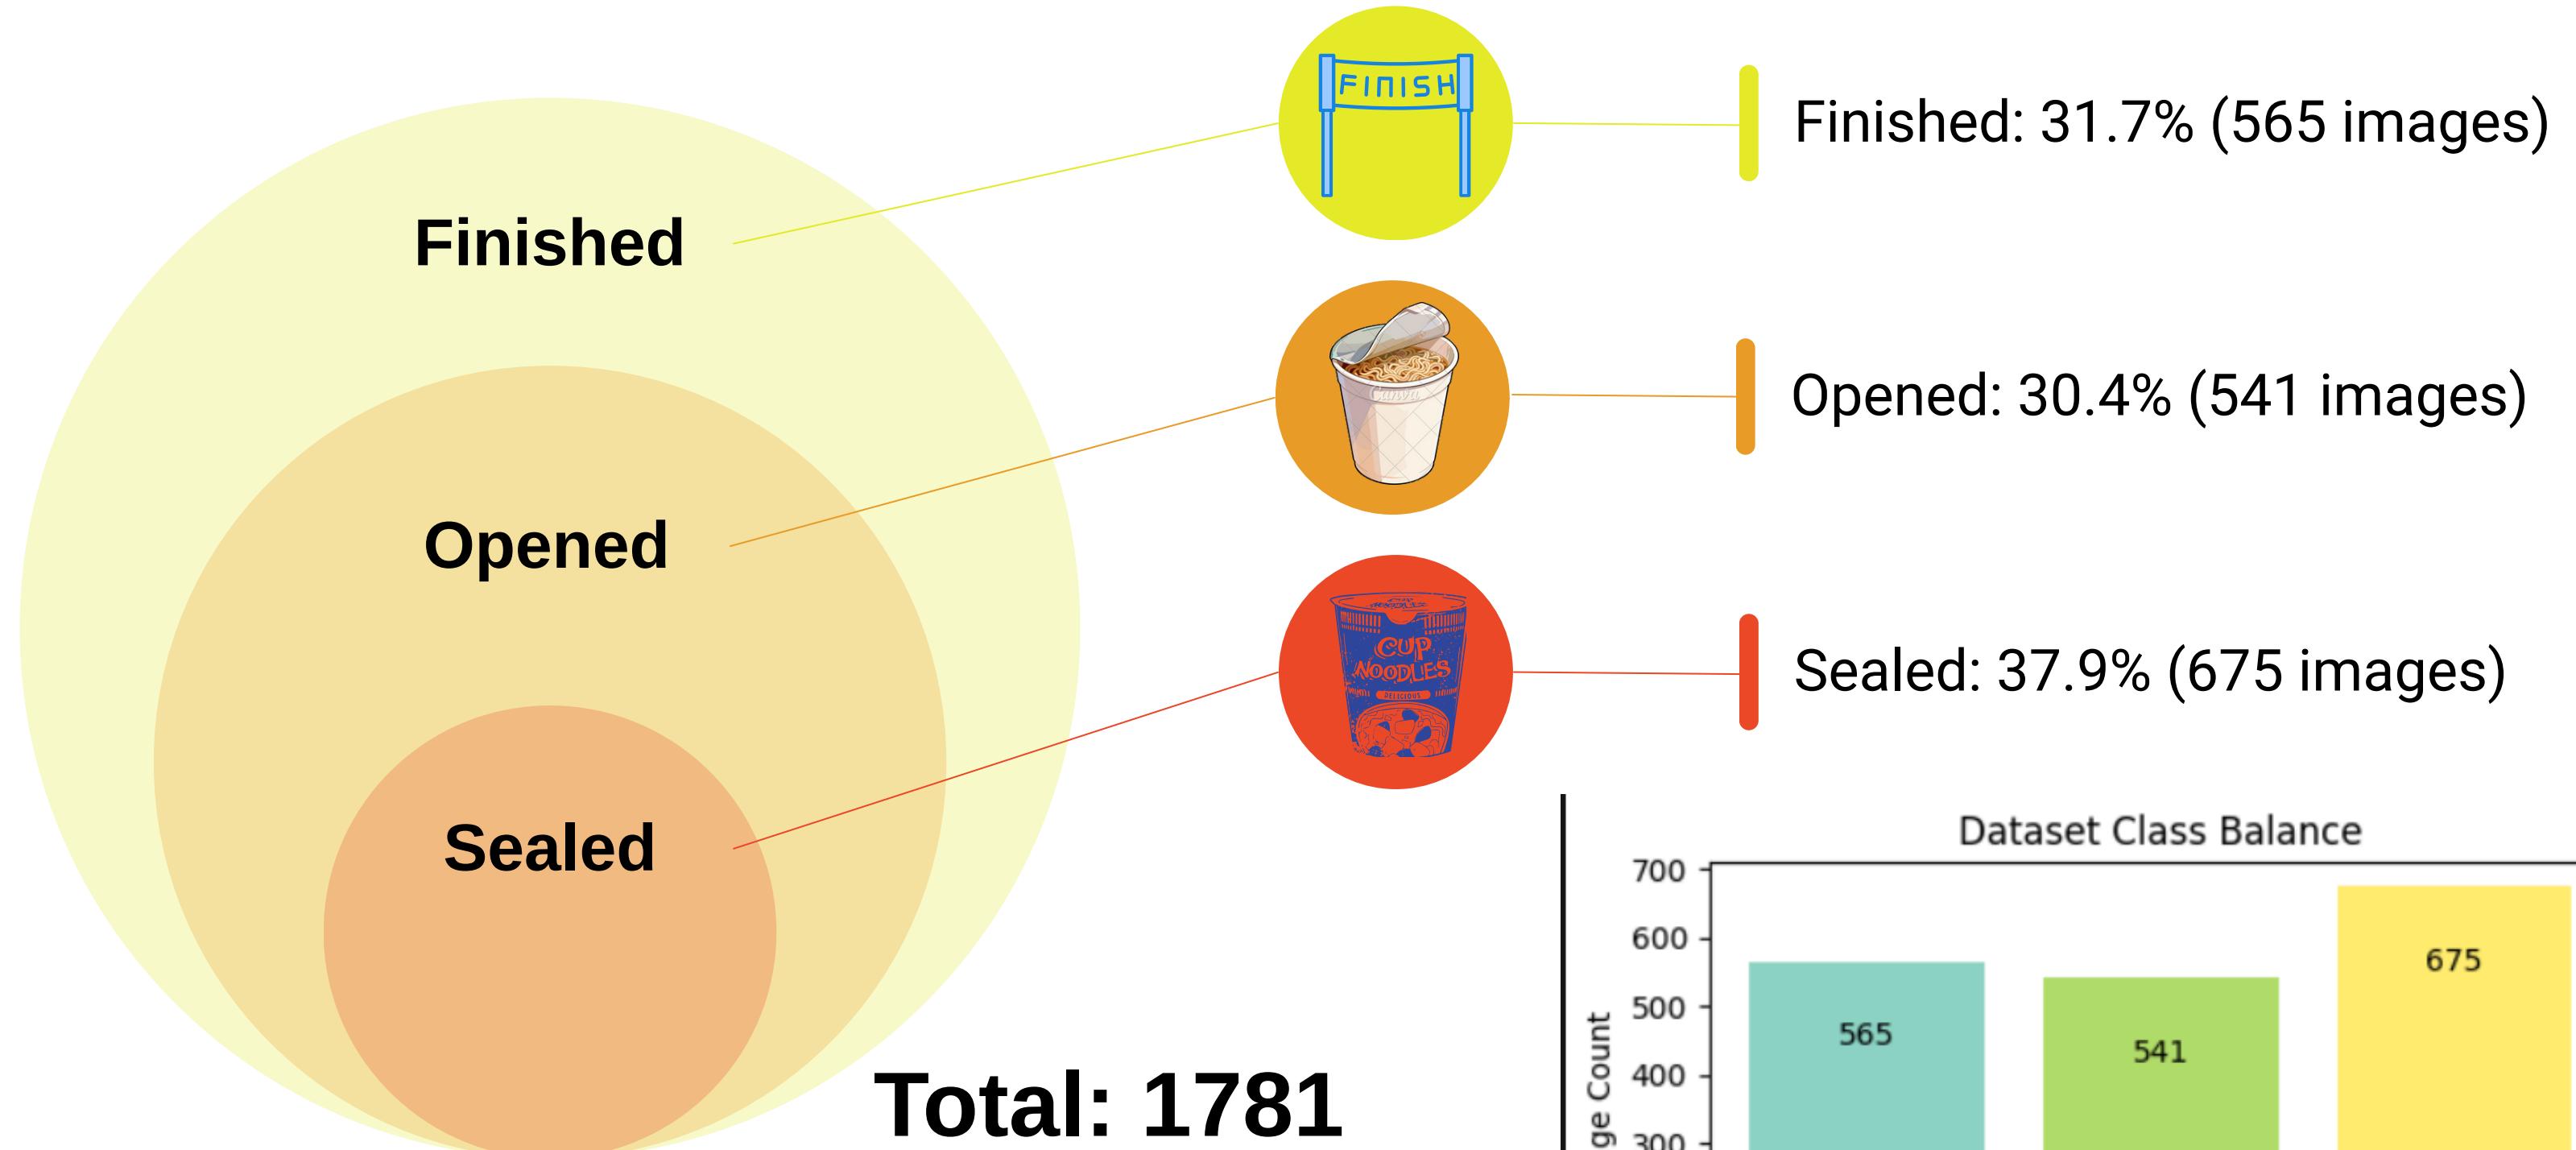
Video to Image Extraction

- Video files were converted into image frames using OpenCV
- Frames were extracted at 6 frames per second (FPS) to balance:
 - Dataset size
 - Redundancy between frames
 - Ensures enough time between frames so they are more unique
- Every extracted frame from a video was classified into:
 - Finished
 - Opened
 - Sealed

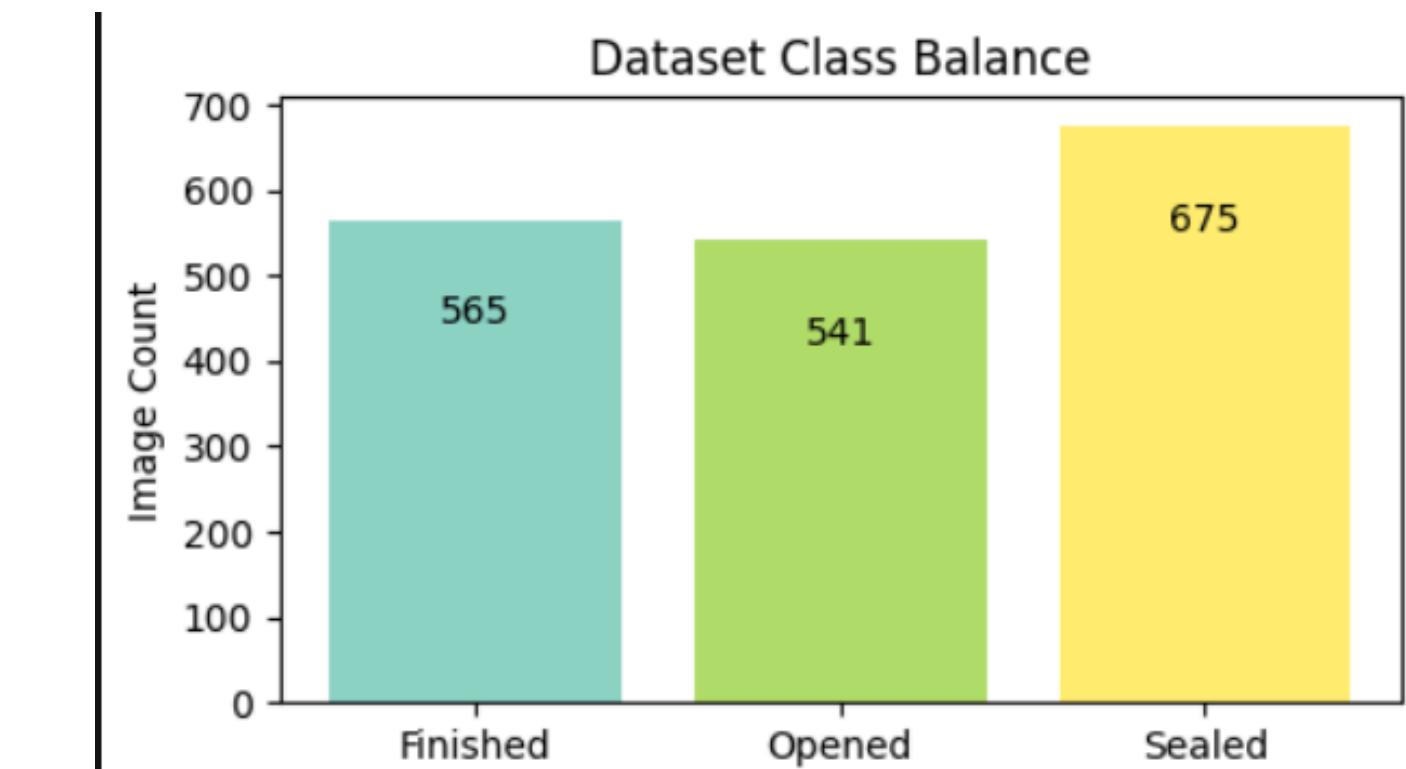
Why this approach was used?

- Increases dataset size without needing more physical samples.
- Captures different angles, lighting conditions, and motion blur.
- Improves model robustness to real-world camera input.
- Reduces overfitting by increasing data diversity.

Dataset Distribution



- Class imbalance handled using class weighting
- Prevents bias towards majority class



Hyperparameters

```
HYPER_PARAMS = {
    'IMG_SIZE': (299, 299), # Max: 299x299 for InceptionV3 but 224x224 also good
    'BATCH_SIZE': 24, # ~24 for Laptop, ~48 for PC
    'EPOCHS': 20, # 5 is ok, 12-18 is when it will early stop on its own
    'DATA_SPLIT': (80, 10, 10), # probably should not change
    'DF_SPLIT_SEED': 42, # controlled seed = replicatable outcomes & fair testing for hyperparam tuning
    'LEARNING_RATE': 1e-3,
    'DROPOUT_RATE': 0.3,
    'DENSE_UNITS': 128,
    'REDUCE_LR_PATIENCE': 2, # x epochs until it reduces Learning rate to prevent bouncing
    'EARLY_STOPPING_PATIENCE': 3, # x epochs until it early stops
    'DISTRIBUTE_CLASS_WEIGHTS': True, # Default: True, will alleviate uneven class image distribution
}
```

Balancing Dataset Class Weights

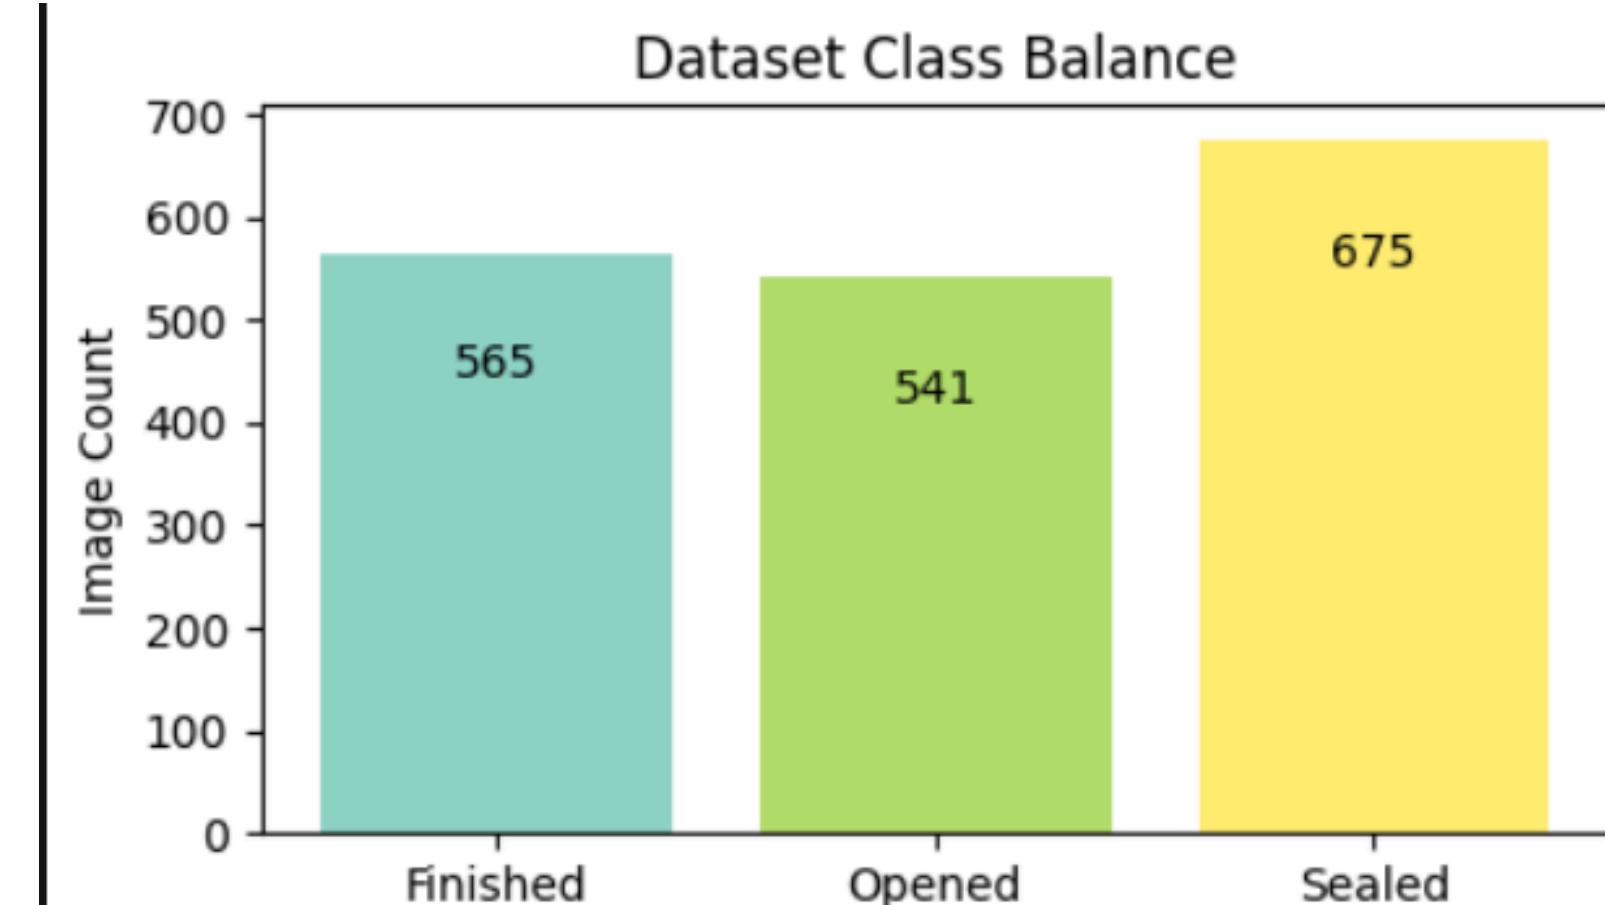
```
print("\n✿ Balancing class weights...")
class_weights = compute_class_weight(
    'balanced', classes=np.unique(train_gen.classes), y=train_gen.classes
)
class_indices = train_gen.class_indices
class_weight_dict = {class_indices[class_name]: weight
                     for class_name, weight in zip(class_indices.keys(), class_weights)}
```

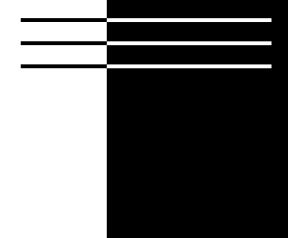
Dataset Class Balance

- Dataset balance is checked before training
- Based on the number of images per class, weights are allocated to each class.
- For example (refer to the bar graph on the right):
 - **Sealed** will have a lower weightage having the **most** images (675)
 - **Opened** will have a higher weightage having the **least** images (541)
- This ensures fairness across classes & hence a better model overall.

`compute_class_weight()` from `sklearn` is used.
`class_weights_dict` is passed into the `model.fit()`

```
history = model.fit(
    train_gen,
    epochs=HYPER_PARAMS['EPOCHS'],
    validation_data=val_gen,
    class_weight=class_weight_dict,
    callbacks=callbacks,
    verbose=1 # prints the cool progress bar
)
```





Data Preprocessing



```
# returns the 3 respective image generators
def build_image_generators():
    img_size = HYPER_PARAMS["IMG_SIZE"]
    batch_size = HYPER_PARAMS["BATCH_SIZE"]
    train_df, val_df, test_df = build_df_splits(data_dir)

    train_datagen = ImageDataGenerator(
        rescale=1./255, # normalizes pixel values
        rotation_range=20,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True # 5 methods of data augmentation
    )
    val_test_datagen = ImageDataGenerator(rescale=1./255) # only normalize pixel values, no data aug
```

Image Preprocessing

- Pixel normalization:
 - $\text{rescale} = 1 / 255$
- Converts pixel values from $0-255 \rightarrow 0-1$
- Ensures consistent input across all images

Data Augmentation

Applied only to training data:

- Rotation
- Width & height shift
- Zoom
- Horizontal flip

Purpose:

- Simulate real-world variations
- Reduce overfitting
- Improve robustness

Dataset Expansion/Splitting

```
# Returns 3 shuffled data frames of the 80/10/10 data split
def build_df_splits(data_dir):
    seed = 42 # seed only controls the df splits
    random.seed(seed)
    np.random.seed(seed)
    tf.random.set_seed(seed)
    files = []
    labels = []

    valid_ext = (".jpg", ".jpeg", ".png")

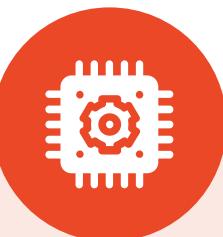
    for cls in class_names:
        cls_dir = os.path.join(data_dir, cls)
        if not os.path.isdir(cls_dir):
            continue

        for f in os.listdir(cls_dir):
            if f.lower().endswith(valid_ext):
                files.append(os.path.join(cls_dir, f)) # gets all filepaths of images
                labels.append(cls) # gets all labels/class names of images

    df = pd.DataFrame({"filename": files, "class": labels})
    df = df.sample(frac=1, random_state=seed).reset_index(drop=True) # randomises its order

    n = len(df) # total number of images
    train_end = int(0.8 * n)
    val_end = int(0.9 * n)

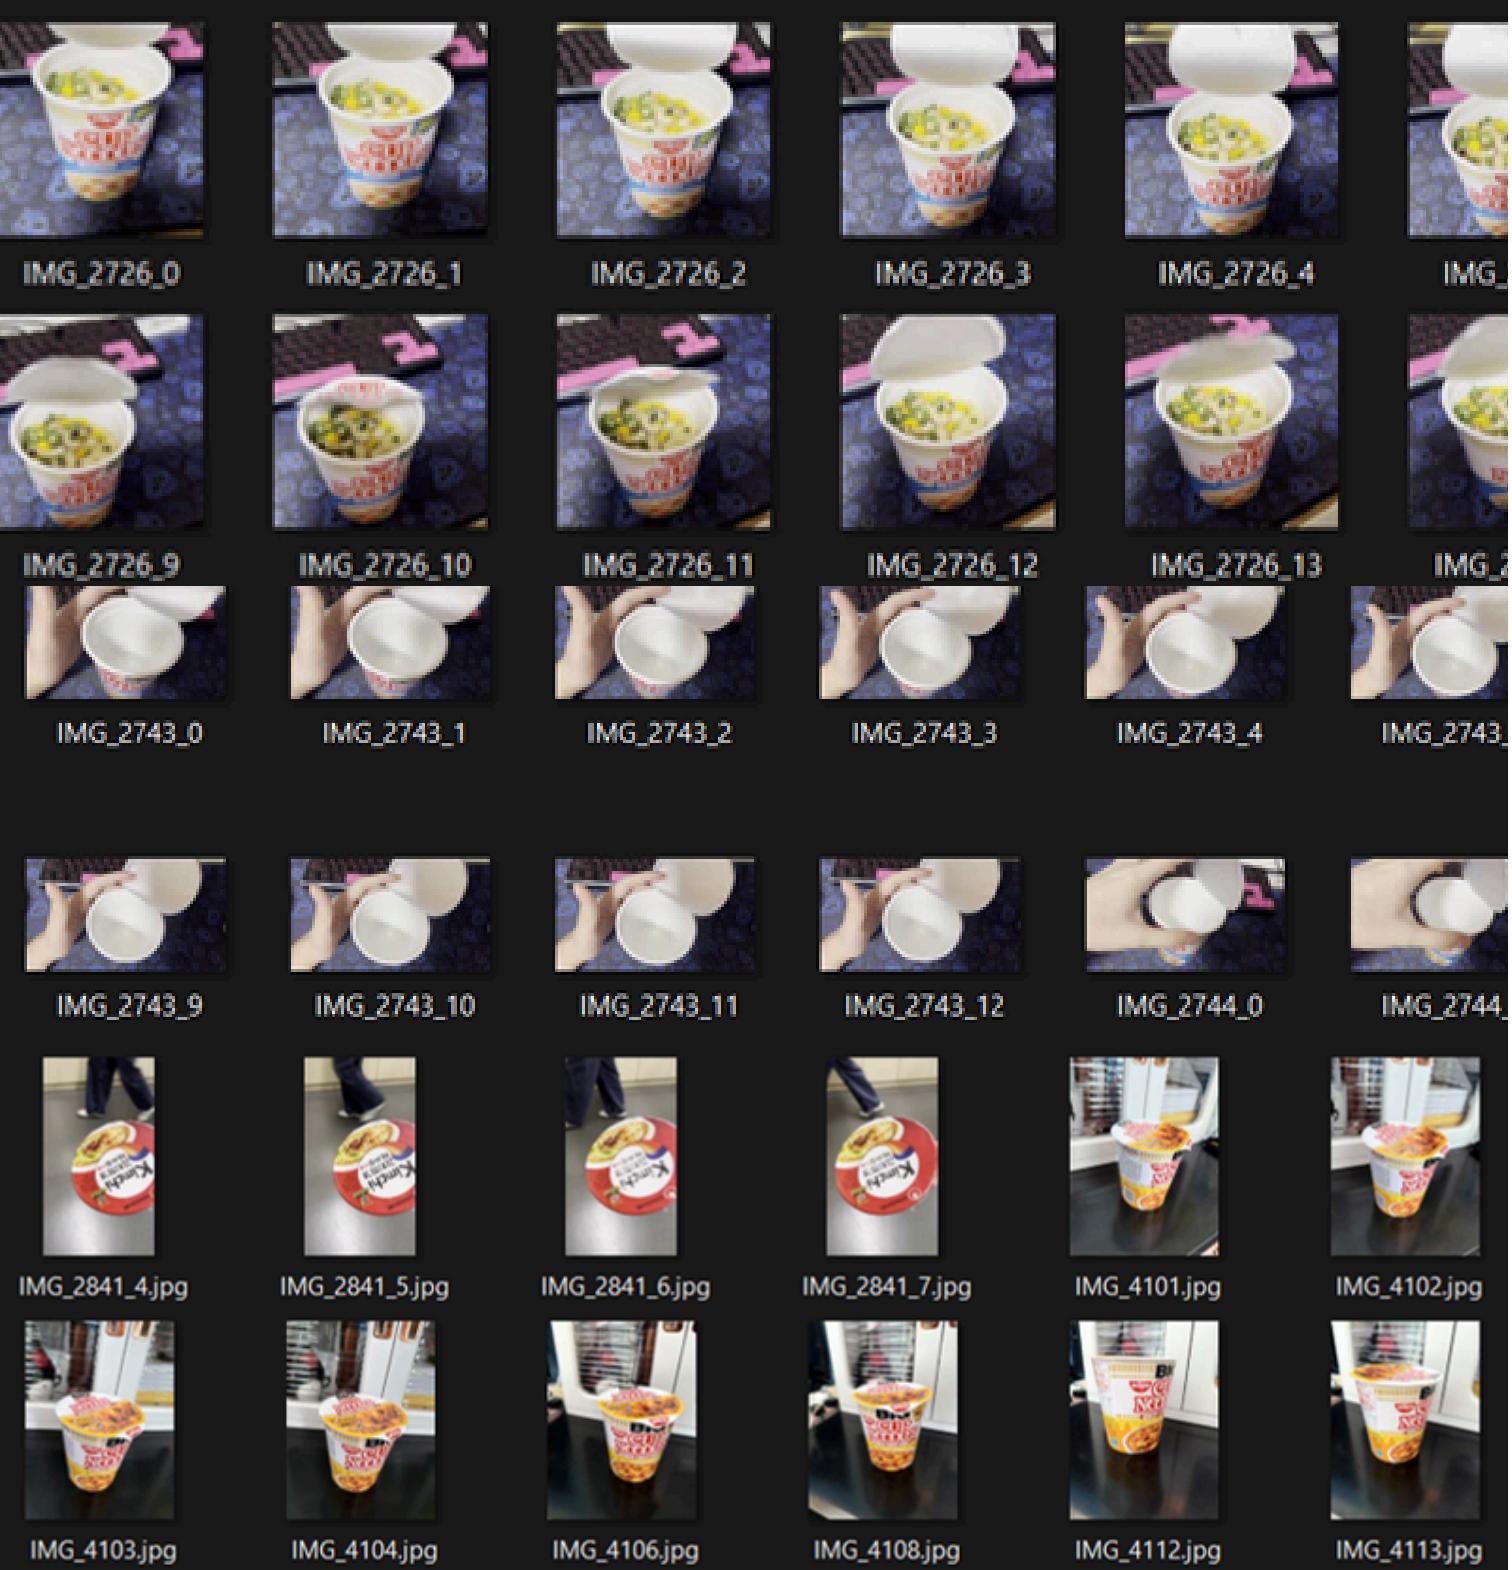
    train_df = df[:train_end] # 80% goes to training dataframe
    val_df = df[train_end:val_end] # 10% goes to validation dataframe
    test_df = df[val_end:] # 10% goes to testing dataframe
```



Data Splitting Strategy (80 / 10 / 10)

- Dataset is split into:
 - 80% Training
 - 10% Validation
 - 10% Testing
- Split is:
 - Randomized
 - Reproducible using a fixed seed
- Prevents data leakage between sets
- Only .jpg, .jpeg, .png used

Model Preparation



Model Preparation

- All images are resized to 299×299
 - Reason:
 - InceptionV3 is pretrained on ImageNet using 299×299 images
 - Preserves more spatial detail compared to smaller inputs
 - Our dataset consists of images of many different sizes
 - Trade-offs:
 - Higher validation accuracy (97.7% at convergence)
 - BUT Longer training time
- Chosen input size: 299×299 for improved accuracy

299x299 (Early stop, same seed) → 13min

✓ Test accuracy: 0.9832

224x224 (Early stop, same seed) → 10min

✓ Test accuracy: 0.9608

Model Architecture

```
def build_model():
    base_model = InceptionV3(input_shape=(*HYPER_PARAMS["IMG_SIZE"], 3), include_top=False, weights='imagenet')
    base_model.trainable = False
    x = base_model.output
    x = GlobalAveragePooling2D()(x)
    x = Dropout(0.3)(x)
    x = Dense(128)(x)
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = Dropout(0.3)(x)
    predictions = Dense(len(class_names), activation='softmax')(x)

    model = Model(inputs=base_model.input, outputs=predictions)
    return model
```

Model Architecture

- **Base model:** InceptionV3
 - Pretrained on ImageNet
- **Feature extractor layers are:**
 - Frozen
 - Not retrained
- Custom classification head added on top

Why InceptionV3?

- Strong feature extraction
- Good performance on real-world images
- Faster convergence
- Higher val_acc than other models (we tried)

Custom Classification Head

Added layers:

- Global Average Pooling
- Dense layer (128 units)
- Batch Normalisation
- ReLU activation
- Dropout (0.3)
- Final Softmax layer (3 classes)

Benefits:

- GAP: reduces parameters
- BatchNorm: stabilizes training
- Dropout: prevents overfitting
- Softmax: outputs probabilities

Model Callbacks

```
callbacks = [  
    EarlyStopping(patience=3, restore_best_weights=True, monitor='val_accuracy'),  
    ReduceLROnPlateau(factor=0.2, patience=2, min_lr=1e-7) # slows learning rate if accuracy plateaus  
]
```

Early Stopping

Stops the model early after **3 epochs** (AKA *patience*) of no improvements in validation accuracy.

Helps save time when the model has possibly reached its minima (converged).

Reduce LR on Plateau

(LR = Learning Rate)

Lowers learning rate of model after **2 epochs** (AKA *patience*) of no improvements in validation accuracy.

Helps the model escape local minima, where it bounces near the minima because learning rate is too high.

Image Size Experiment

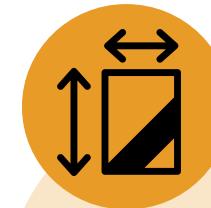
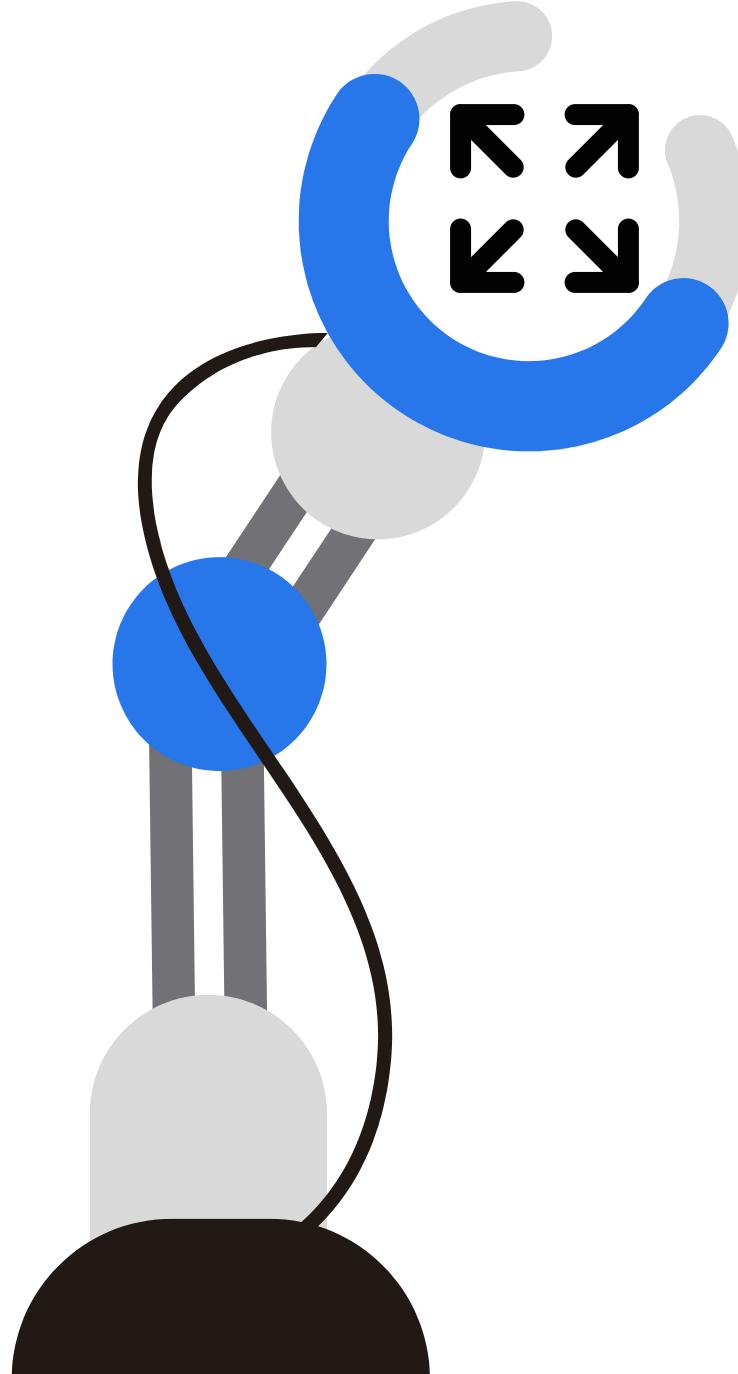


Image Size Experiment

- Initial experiments used smaller image sizes
- Later experiments increased image size
- Observation:
 - Larger image size improved validation accuracy (12.7%)
 - Training time increased (31.3%)

Conclusion:

Increasing image size improves accuracy when using transfer learning models like InceptionV3

```
# ===== CONFIG =====  
  
IMG_SIZE = (75, 75)  
BATCH_SIZE = 32  
EPOCHS = 8  
SEED = 42
```

Smaller image size (75x75)

accuracy: 0.7959

```
HYPER_PARAMS = {  
    'IMG_SIZE': (299, 299), # Max: 299  
    'BATCH_SIZE': 48, # ~24 for Laptop  
    'EPOCHS': 1, # 12-18 is when it works  
    'DISTRIBUTE_CLASS_WEIGHTS': True  
}
```

Larger image size (299x299)

accuracy: 0.8968

Experiment Comparison



Experiment Comparison

- Average training time increased:
 - +45.8%
- Final validation accuracy improved:
 - +12.7%
- Trade-off:
 - Slower training
 - Higher accuracy

Conclusion:

The improvement in accuracy is worth the increased training time

```
Epoch 1/20
27/27 122s 4s/step - accuracy: 0.5990 - loss: 1.2552 - val_accuracy: 0.8101 - val_loss: 0.4866
Epoch 2/20
27/27 111s 4s/step - accuracy: 0.7239 - loss: 0.8375 - val_accuracy: 0.7475 - val_loss: 0.5711
Epoch 3/20
27/27 121s 5s/step - accuracy: 0.7743 - loss: 0.6203 - val_accuracy: 0.8255 - val_loss: 0.4284
Epoch 4/20
27/27 120s 4s/step - accuracy: 0.7911 - loss: 0.5140 - val_accuracy: 0.8441 - val_loss: 0.3921
Epoch 5/20
27/27 120s 5s/step - accuracy: 0.7959 - loss: 0.5301 - val_accuracy: 0.8661 - val_loss: 0.3444
```

IMG_SIZE: 75x75

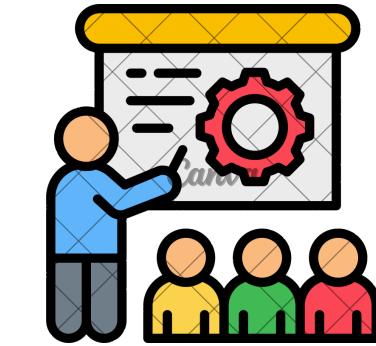
accuracy: 0.7959

```
Epoch 1/5
35/35 192s 5s/step - accuracy: 0.6466 - loss: 1.0876 - val_accuracy: 0.8666 - val_loss: 0.3747
Epoch 2/5
35/35 170s 5s/step - accuracy: 0.8018 - loss: 0.4843 - val_accuracy: 0.9058 - val_loss: 0.2629
Epoch 3/5
35/35 166s 5s/step - accuracy: 0.8457 - loss: 0.3841 - val_accuracy: 0.8741 - val_loss: 0.3043
Epoch 4/5
35/35 178s 5s/step - accuracy: 0.8685 - loss: 0.3583 - val_accuracy: 0.9374 - val_loss: 0.1820
Epoch 5/5
35/35 160s 5s/step - accuracy: 0.8968 - loss: 0.3069 - val_accuracy: 0.9466 - val_loss: 0.1674
```

IMG_SIZE: 224x224

accuracy: 0.8968

Training Configuration

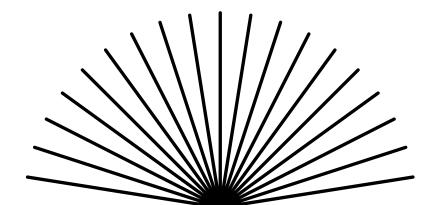


Training Configuration

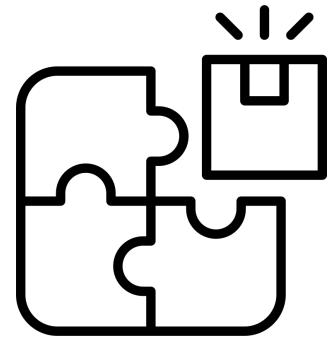
- Optimizer: Adam
- Loss function: Categorical Cross-Entropy
- Learning rate adjusted dynamically
- Callbacks:
 - EarlyStopping
 - ReduceLROnPlateau
- Class weights applied to balance dataset

```
# 3. Build the model's layers
print("\n\tBuilding model...")
model = build_model()
model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=1e-3),
    loss='categorical_crossentropy',
    metrics=['accuracy'])
)

# 4. Training the model
print("\n\tTraining...")
callbacks = [
    EarlyStopping(patience=3, restore_best_weights=True, monitor='val_accuracy'), # patience means wait x epochs before doing smth
    ReduceLROnPlateau(factor=0.2, patience=2, min_lr=1e-7) # slows Learning rate to prevent it from bouncing
]
start_time = time.time()
history = model.fit(
    train_gen,
    epochs=HYPER_PARAMS['EPOCHS'],
    validation_data=val_gen,
    class_weight=class_weight_dict,
    callbacks=callbacks,
    verbose=1 # prints the cool progress bar
)
training_time = time.time() - start_time
model accuracy = model.evaluate(test_gen)[1]
```



Overfitting Prevention Techniques



```
# returns the 3 respective image generators
def build_image_generators():
    img_size = HYPER_PARAMS["IMG_SIZE"]
    batch_size = HYPER_PARAMS["BATCH_SIZE"]
    train_df, val_df, test_df = build_df_splits(data_dir)

    train_datagen = ImageDataGenerator(
        rescale=1./255, # normalizes pixel values
        rotation_range=20,
        width_shift_range=0.2,
        height_shift_range=0.2,
        zoom_range=0.2,
        horizontal_flip=True # 5 methods of data augmentation
    )
    val_test_datagen = ImageDataGenerator(rescale=1./255) # only normalize pixel values, no data aug
```

Techniques Used

- Data augmentation
- Dropout layers
- Batch normalization
- Early stopping
- Learning rate reduction
- Class weighting

Effectiveness

- Reduced training-validation gap
- Improved stability

```
# 4. Training the model
print("\n\x1b[31m Training...")
callbacks = [
    EarlyStopping(patience=3, restore_best_weights=True, monitor='val_accuracy'), # patience means wait x epochs before doing smth
    ReduceLROnPlateau(factor=0.2, patience=2, min_lr=1e-7) # slows Learning rate to prevent it from bouncing
]
```

Dataset Class Balance

Dataset Class Balance

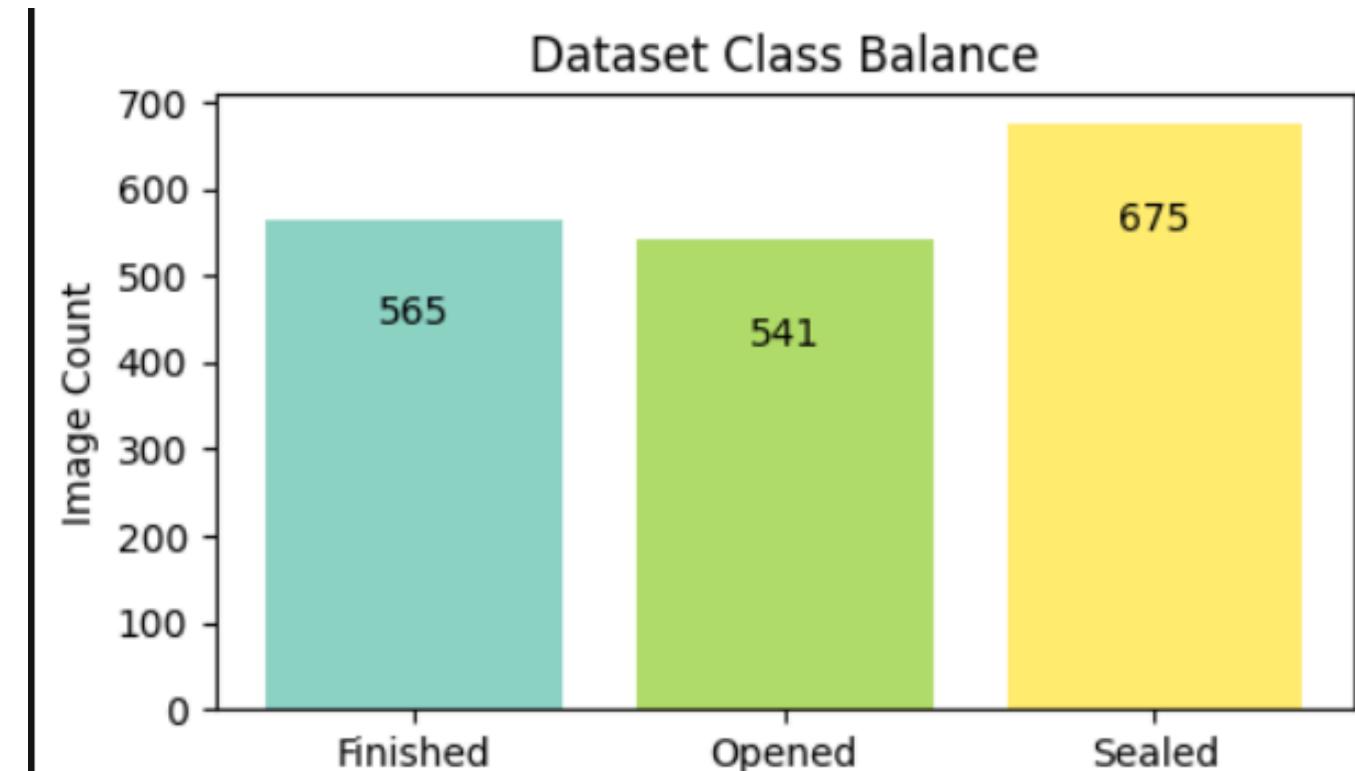
- Dataset balance is checked before training
- Number of images per class is displayed
- Bar chart used to visualise class distribution
- Class imbalance handled using:
 - Computed class weights
 - Prevents bias towards dominant classes
- Visualization created using Matplotlib bar charts

```
def check_dataset_balance():
    counts = {}
    for class_dir in sorted(class_names):
        class_path = os.path.join(data_dir, class_dir)
        if os.path.isdir(class_path):
            counts[class_dir] = len([f for f in os.listdir(class_path)
                                   if f.lower().endswith('.jpg', '.png', '.jpeg')])

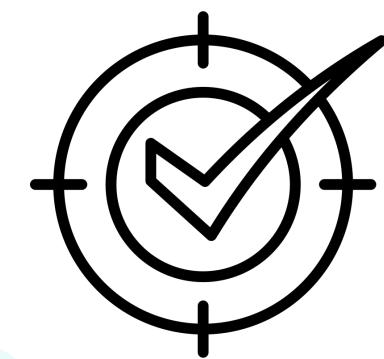
    total = sum(counts.values())
    print("Dataset Balance:")
    print(counts)
    print("Ratios:", {k: f"{v/total*100:.1f}%" for k,v in counts.items()})

    plt.figure(figsize=(5, 4))
    colors = plt.cm.Set3(np.linspace(0, 1, len(class_names)))
    bars = plt.bar(class_names, counts.values(), color=colors)
    plt.bar_label(bars, padding=-30)
    plt.ylabel('Image Count')
    plt.title('Dataset Class Balance')
    plt.tight_layout()
    plt.show()

check_dataset_balance()
```



Accuracy & Loss Analysis

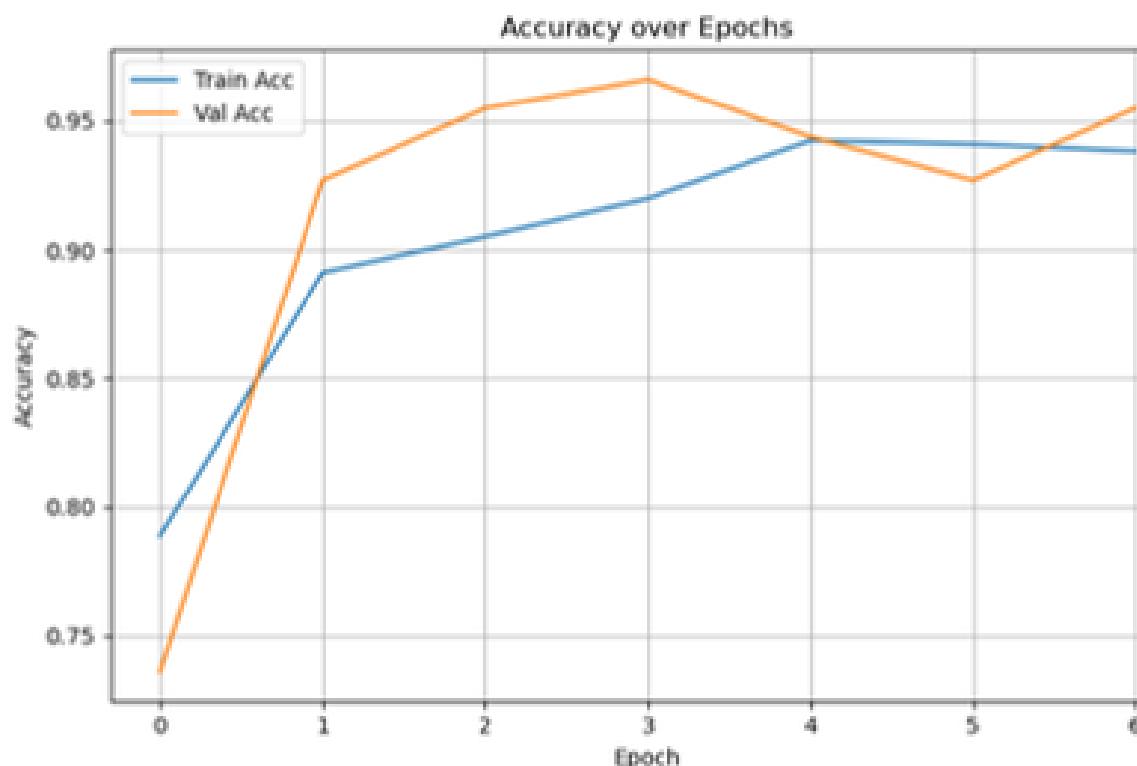


```
print(history.history)
fig, axes = plt.subplots(1, 2, figsize=(15, 5))

# Accuracy
axes[0].plot(history.history['accuracy'], label='Train Acc')
axes[0].plot(history.history['val_accuracy'], label='Val Acc')
axes[0].set_title('Accuracy over Epochs')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Accuracy')
axes[0].legend()
axes[0].grid(True)

# Loss
axes[1].plot(history.history['loss'], label='Train Loss')
axes[1].plot(history.history['val_loss'], label='Val Loss')
axes[1].set_title('Loss over Epochs')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Loss')
axes[1].legend()
axes[1].grid(True)

plt.tight_layout()
plt.show()
plt.close()
```



Why this matters

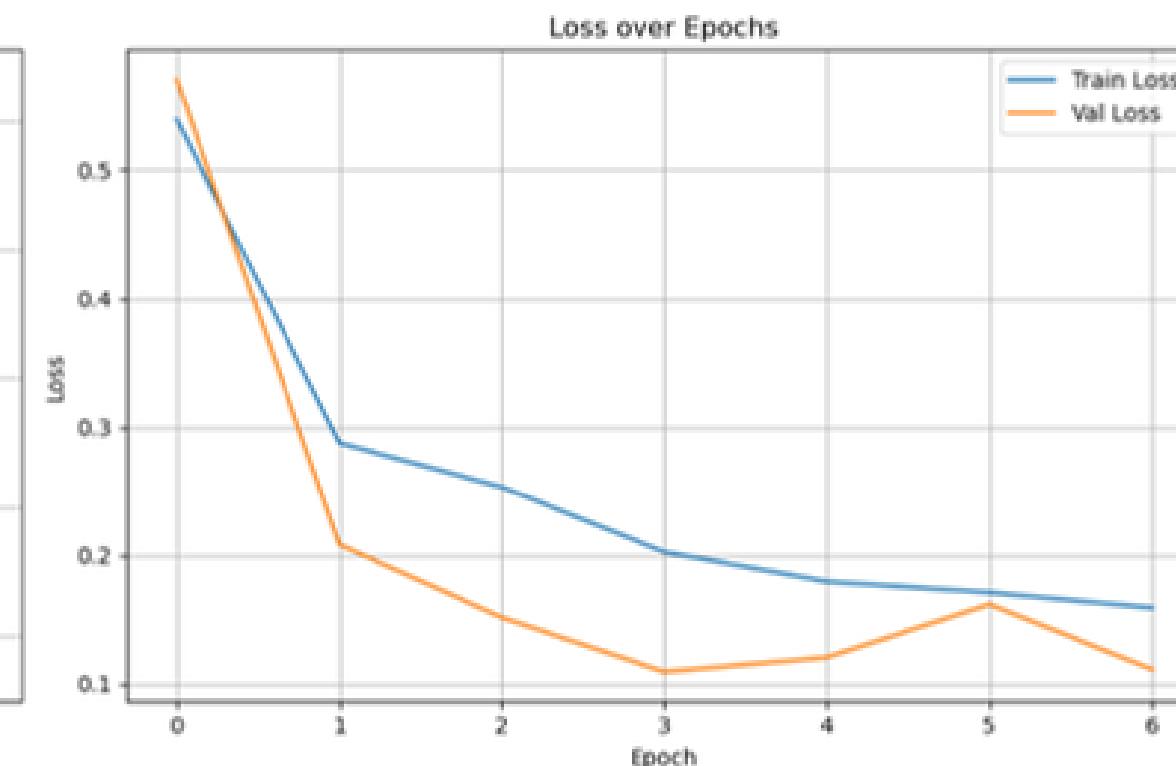
- Detect overfitting
- Check convergence behaviour

What we plotted

- Training Accuracy vs Validation Accuracy
- Training Loss vs Validation Loss

Observation

- Validation curves closely follow training curves
- No severe overfitting



Confusion Matrix

```
if 'y_pred_prob' not in locals(): # skips code if already have the values
    print("Computing predictions...")
    test_gen.reset() # resets generator's pointer to 0
    y_pred_prob = model.predict(test_gen, verbose=0)
    y_pred = np.argmax(y_pred_prob, axis=1) # predictions
    y_true = test_gen.classes # labels

cm = confusion_matrix(y_true, y_pred)
cmdisplay = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=class_names)
cmdisplay.plot(cmap=plt.cm.Blues, xticks_rotation=45)
plt.title('Confusion Matrix (From Test Dataset)')
plt.show()
plt.close()
```

Purpose

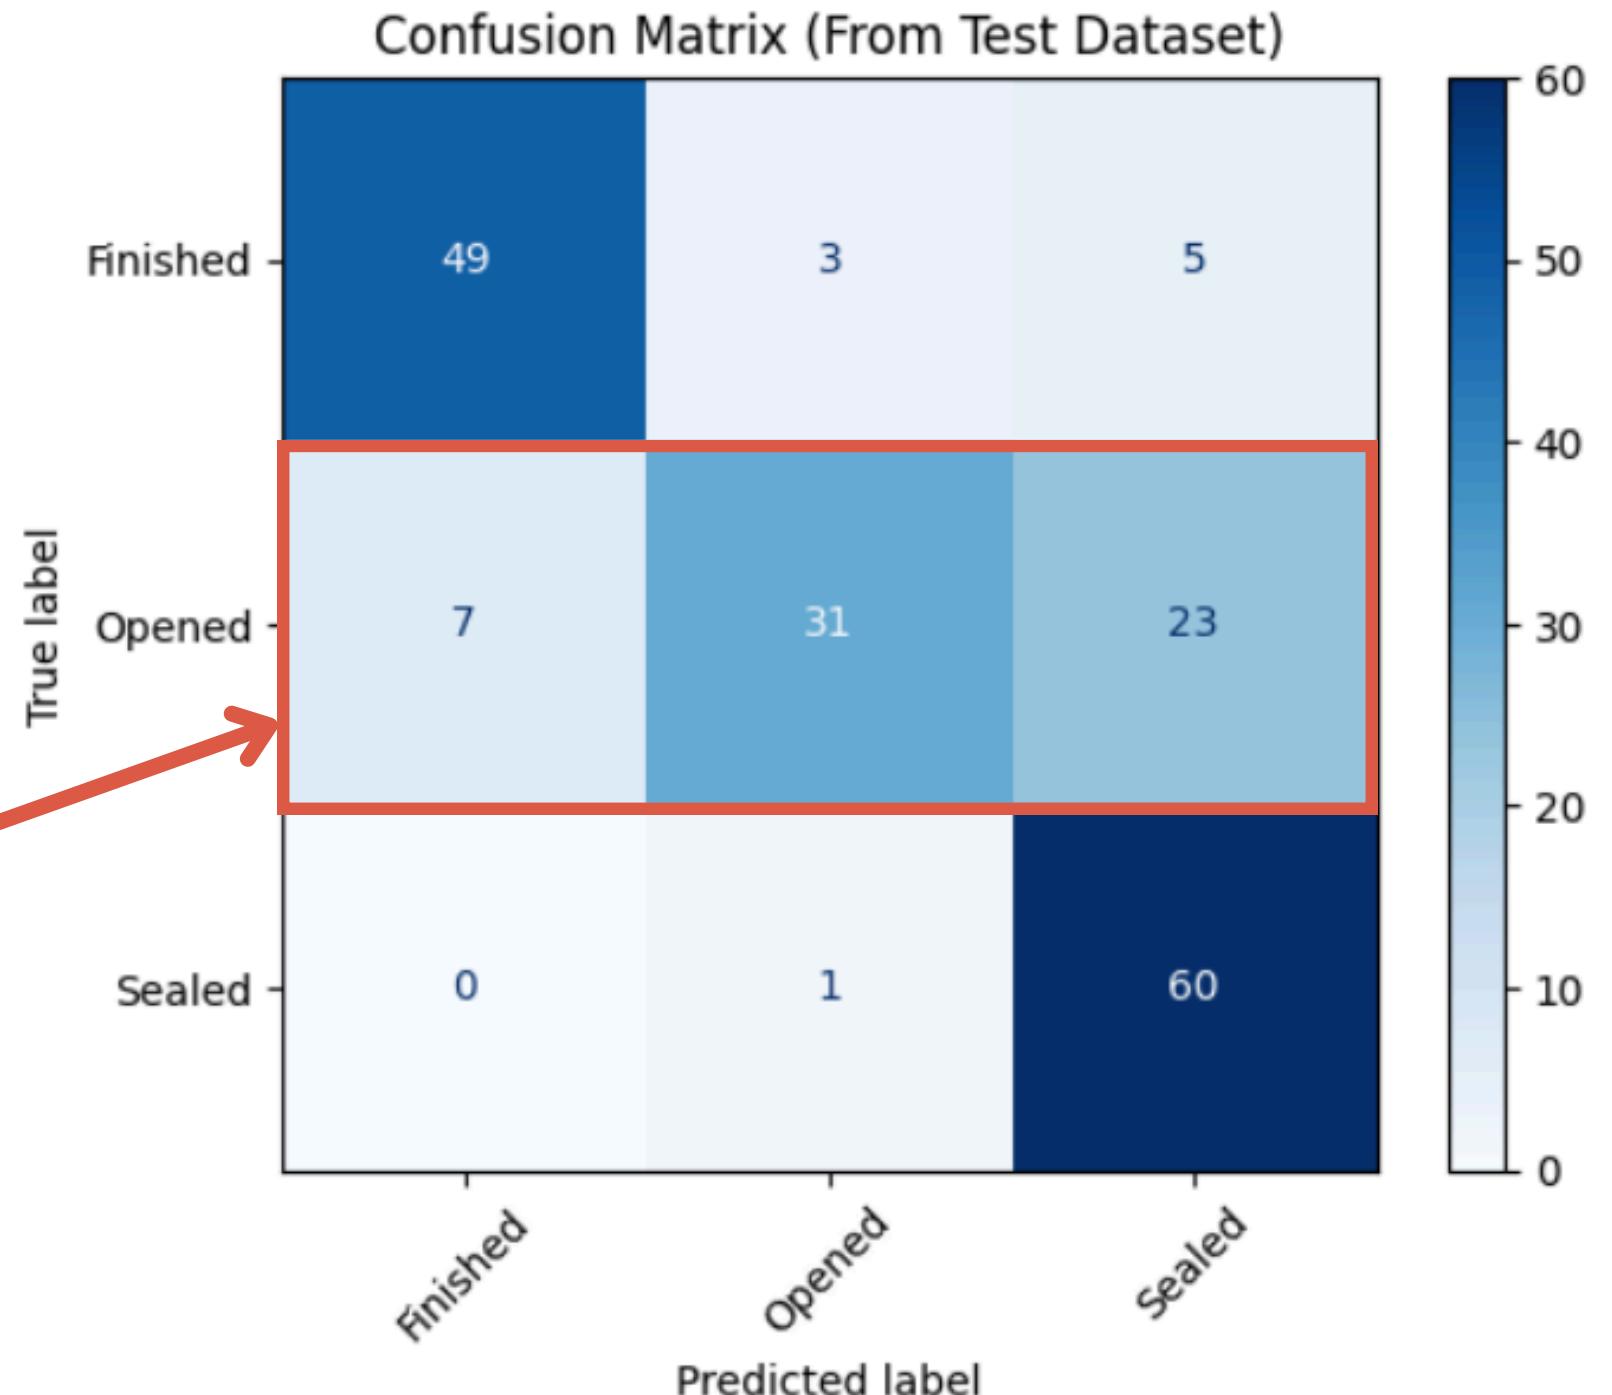
- Identify which classes confuses the model
- Analyse misclassification patterns

Findings

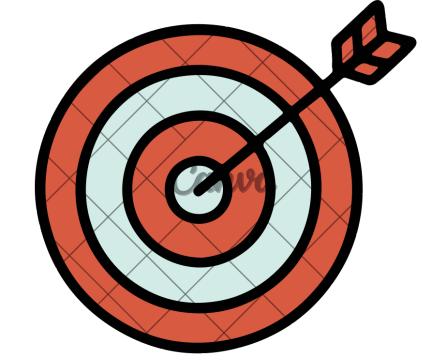
- Most errors occur between:
 - Opened ↔ Sealed
- Finished class highly distinct

Example (refer to the image):

- The **opened** class has the most **wrongly predicted** images (model confused) - so more images may be required in that class for training.



Precision, Recall & F1 Score



```
if 'y_pred_prob' not in locals():
    print("Computing predictions...")
    test_gen.reset()
    y_pred_prob = model.predict(test_gen, verbose=0)
    y_pred = np.argmax(y_pred_prob, axis=1)
    y_true = test_gen.classes

print(classification_report(y_true, y_pred, target_names=class_names))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Finished | 0.88 | 0.86 | 0.87 | 57 |
| Opened | 0.89 | 0.51 | 0.65 | 61 |
| Sealed | 0.68 | 0.98 | 0.81 | 61 |
| accuracy | | | 0.78 | 179 |
| macro avg | 0.81 | 0.78 | 0.77 | 179 |
| weighted avg | 0.81 | 0.78 | 0.77 | 179 |

Why metrics beyond accuracy?

- Accuracy alone hides class-specific errors

Metrics Used

- Precision
- Recall
- F1-Score
- Support

Interpretation

- Balanced precision & recall
- Model not biased towards any class

Example (Image on left)

- Opened has a much higher precision than recall, which means the model is too **careful** in that aspect.
- Sealed has a much higher recall than precision, which means the model is **many false alarms** there.

Misclassified Image Analysis



```
COLUMNS = 4
SORT_BY_DESCENDING_CONFIDENCE = True

if 'y_pred_prob' not in locals(): # skips code if already have the values
    print("Computing predictions...")
    test_gen.reset() # resets generator's pointer to 0
    y_pred_prob = model.predict(test_gen, verbose=0)
    y_pred = np.argmax(y_pred_prob, axis=1) # predictions
    y_true = test_gen.classes # labels

mis_idx = np.where(y_true != y_pred)[0]
confidences = np.max(y_pred_prob[mis_idx], axis=1)
misclassified = test_df.iloc[mis_idx].copy()
misclassified['pred_class'] = [class_names[p] for p in y_pred[mis_idx]]
misclassified['confidence'] = confidences

misclassified = misclassified.sort_values('confidence', ascending=not SORT_BY_DESCENDING_CONFIDENCE) # Sort by highest confidence first

print(f"Found {len(misclassified)} misclassified images")

count = len(misclassified)
rows = (count + COLUMNS - 1) // COLUMNS
```



Angle too low to see inside contents



Seal covers too much to be “OPENED”



Poor lighting - Noodles too bright; blends in

Purpose

- Identify difficult samples
- Improve dataset quality

Approach

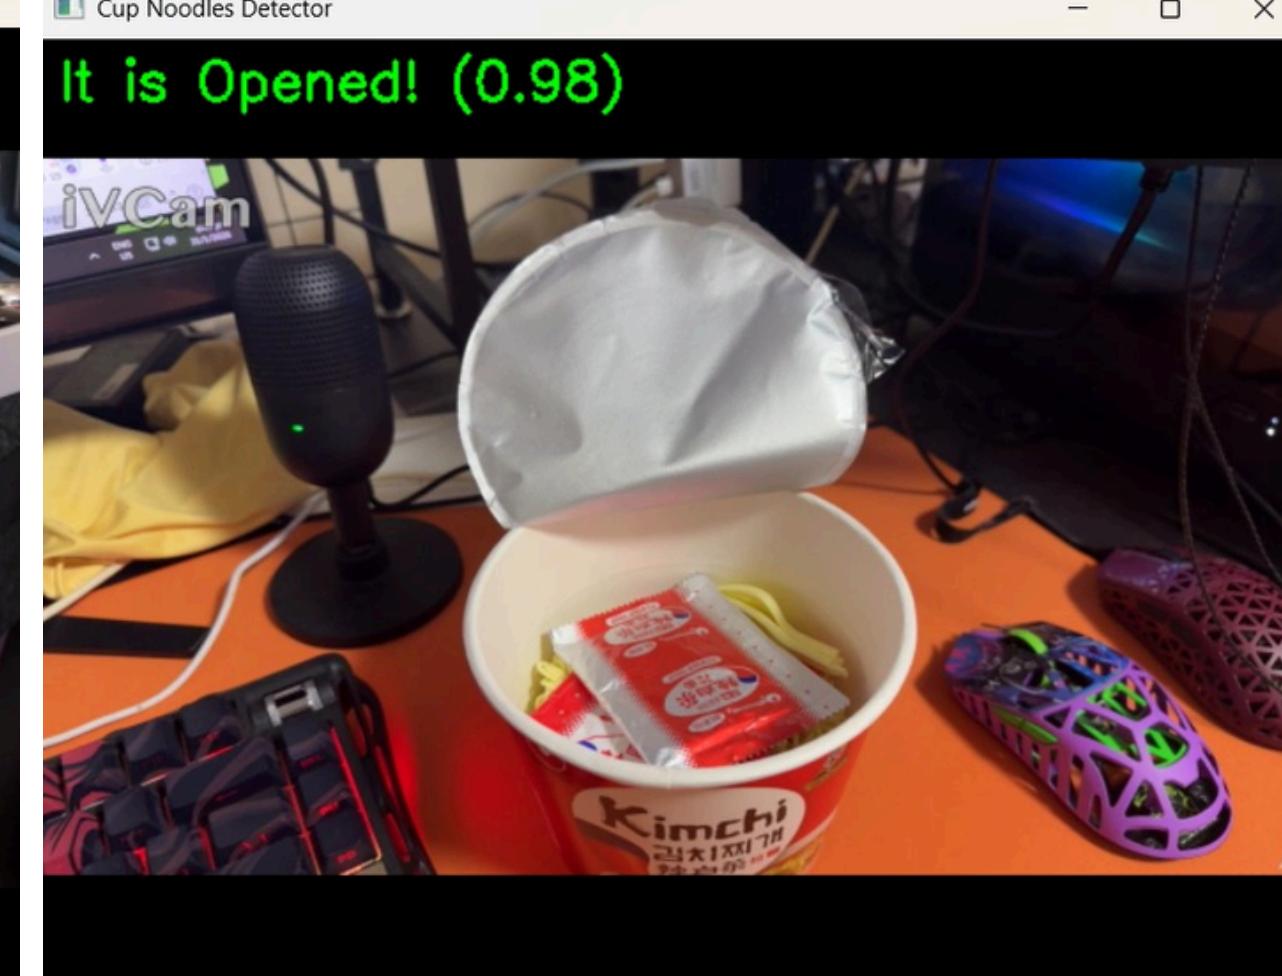
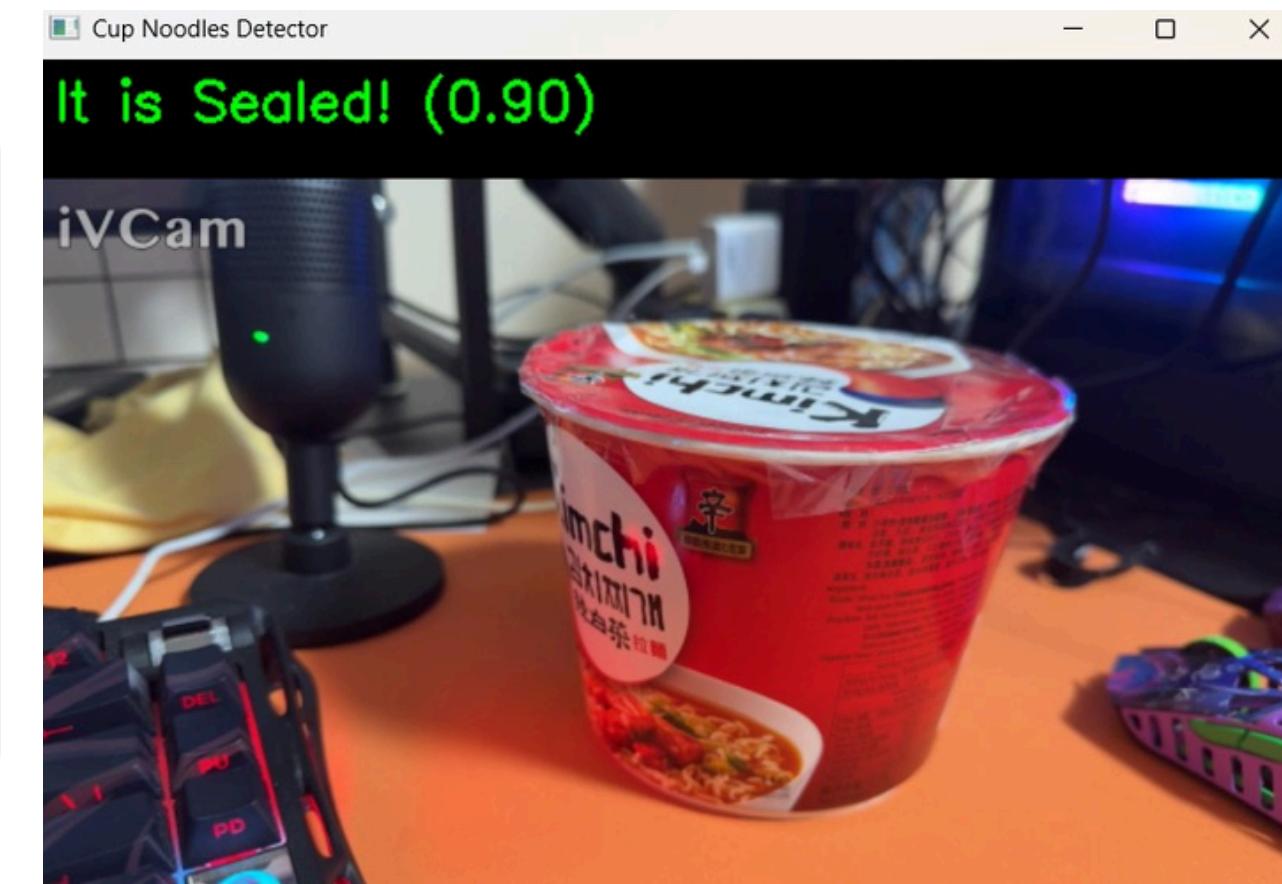
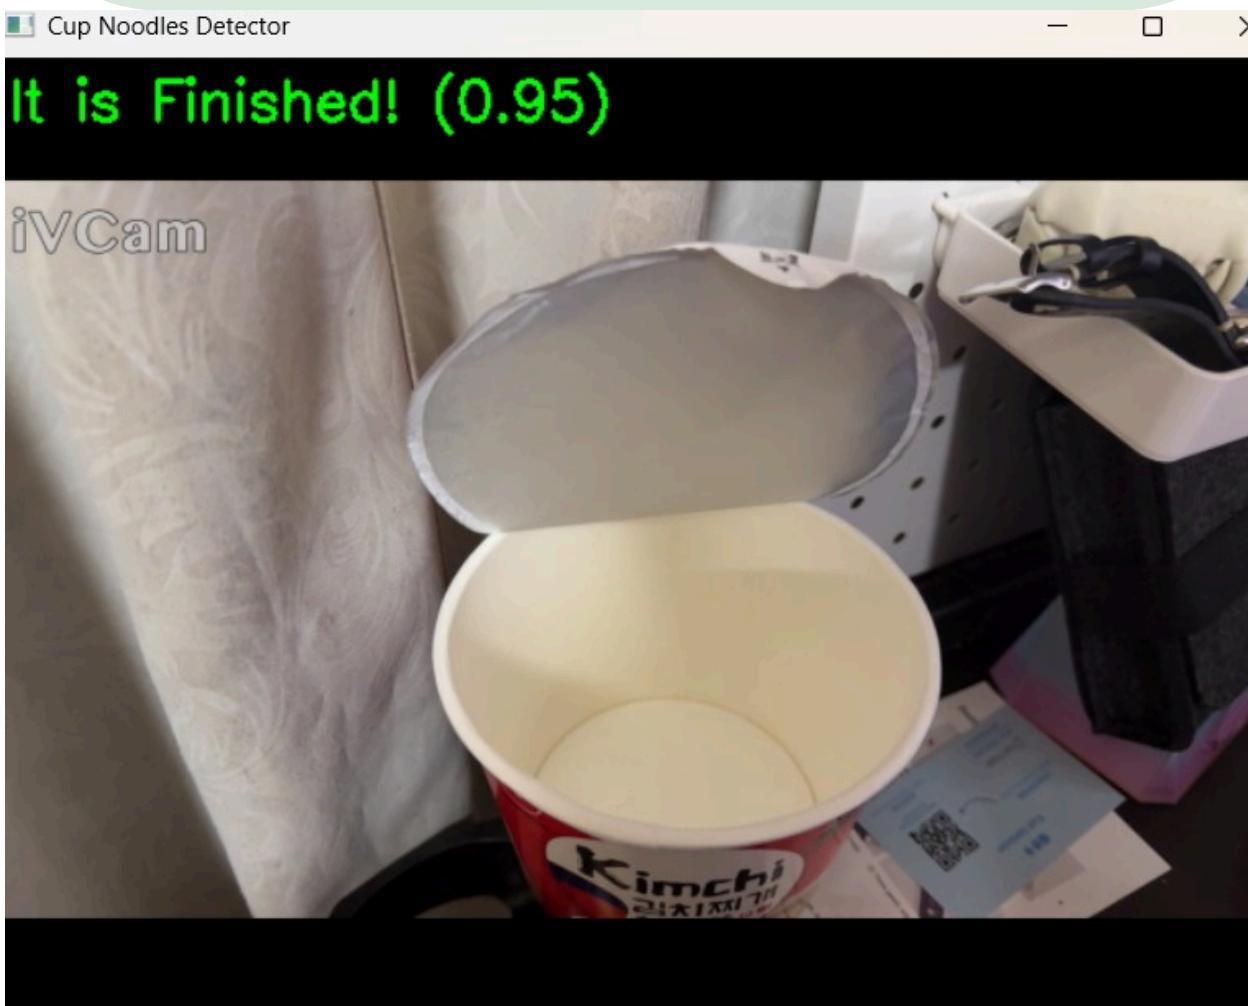
- Extract wrongly predicted images
- Sort by confidence
- Visual inspection

Insights

- Some images:
 - Poor lighting
 - Ambiguous content
 - Confusing angles
- Helps guide future data collection

Live Camera Detection

- Uses OpenCV camera stream
- Each frame is:
 - Preprocessed
 - Passed through the trained model
- Output displayed in real time:
 - Predicted class
 - Confidence score
- Press Q to stop the camera



Training Results



```
📝 Training...
Epoch 1/6
60/60 353s 5s/step - accuracy: 0.7992 - loss: 0.5533 - val_accuracy: 0.8876 - val_loss: 0.3517 - learning_rate: 0.0010
Epoch 2/6
60/60 212s 4s/step - accuracy: 0.9031 - loss: 0.2524 - val_accuracy: 0.9157 - val_loss: 0.2320 - learning_rate: 0.0010
Epoch 3/6
60/60 210s 3s/step - accuracy: 0.9115 - loss: 0.2582 - val_accuracy: 0.9382 - val_loss: 0.1736 - learning_rate: 0.0010
Epoch 4/6
60/60 242s 4s/step - accuracy: 0.9305 - loss: 0.1952 - val_accuracy: 0.9438 - val_loss: 0.1448 - learning_rate: 0.0010
Epoch 5/6
60/60 246s 4s/step - accuracy: 0.9361 - loss: 0.1712 - val_accuracy: 0.9551 - val_loss: 0.1174 - learning_rate: 0.0010
Epoch 6/6
60/60 232s 4s/step - accuracy: 0.9487 - loss: 0.1467 - val_accuracy: 0.9326 - val_loss: 0.1545 - learning_rate: 0.0010
8/8 24s 3s/step - accuracy: 0.9777 - loss: 0.0684
⌚ Total training time: (25.14 minutes)
```

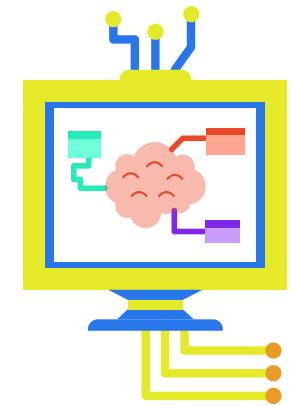
Model Performance – Test Accuracy

- Test Accuracy: 97.7%
- The test dataset contains images not used during training or validation
- Measures the model's ability to generalise to unseen data
- The model correctly classifies about 98 out of every 100 images

This indicates that the model is:

- Reliable
- Not overfitted
- Able to generalise well to new images
- Suitable for real-world image classification tasks

Model Saving Strategy



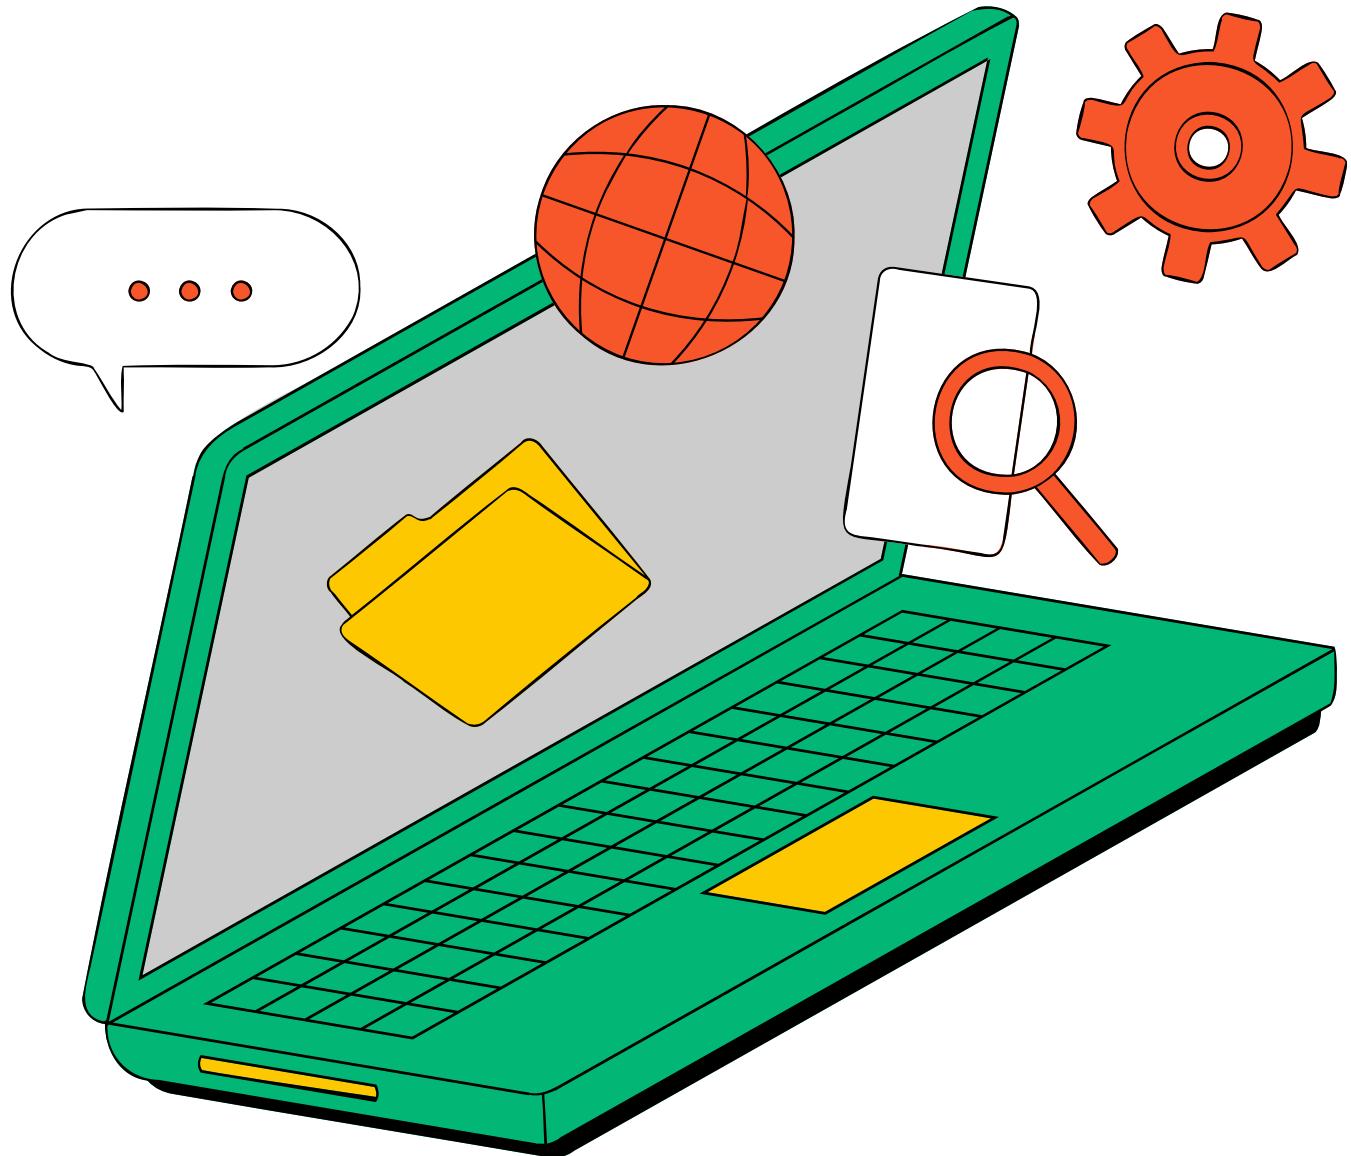
```
# 5. Simple post-training metrics
print(f"⌚ Total training time: ({training_time/60:.2f} minutes)")
print("\n✅ Test accuracy:", model_accuracy)

# 6. Save the model as a .keras file with accuracy & timestamp
timestamp = datetime.datetime.now().strftime("%b%d_%H%M")
filename = f"{int(model_accuracy * 100)}acc_model_{timestamp}.keras"
model.save(f'./models/{filename}')
print(f"💾 Saved: {filename}")
```

Model Saving Strategy

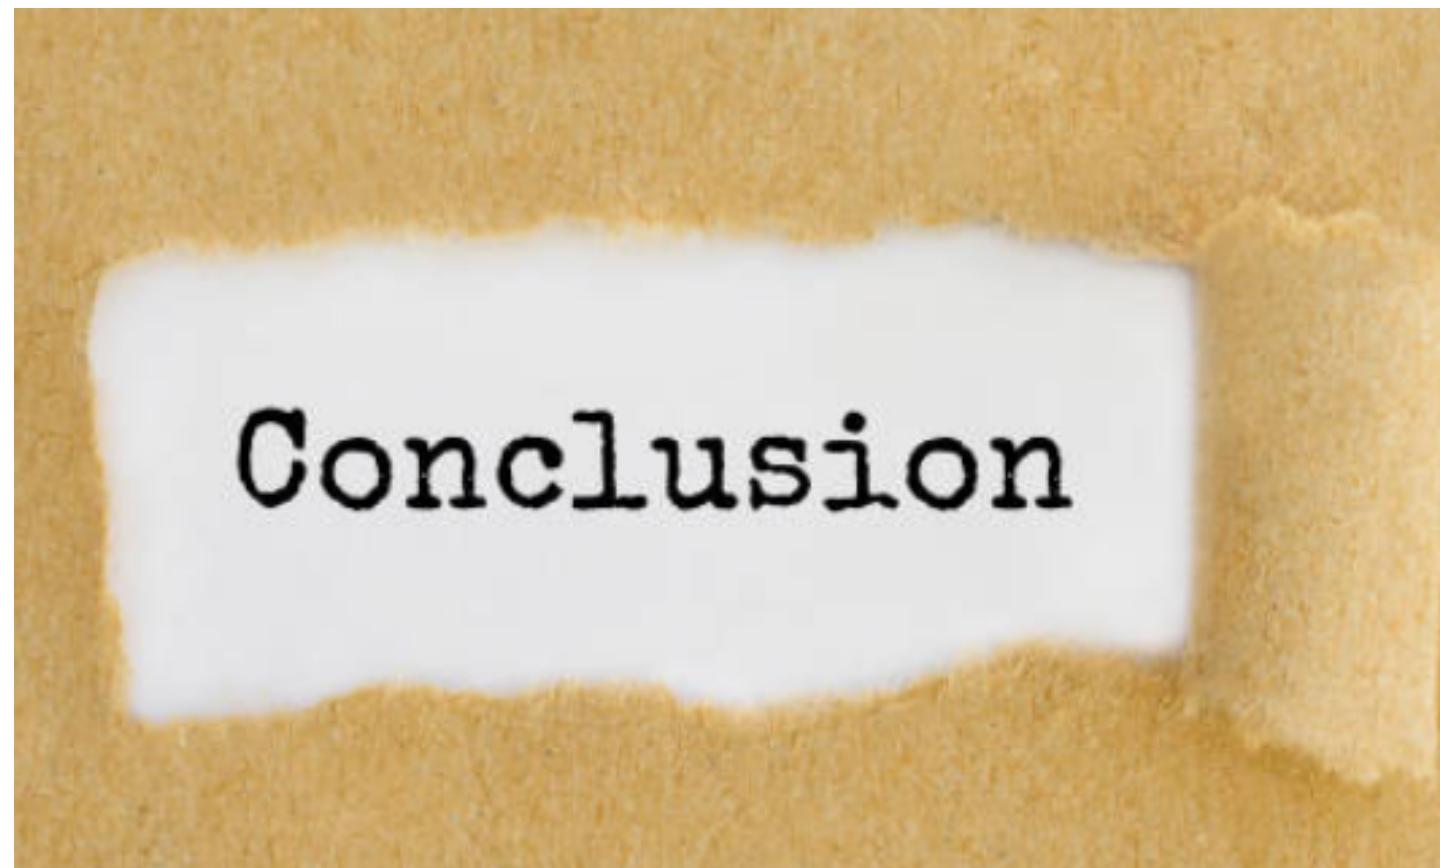
- Model saved in .keras format
- Filename includes:
 - Test accuracy
 - Timestamp
- Allows:
 - Easy version comparison
 - Model reuse without retraining

Improvements & Future Work



- Increase dataset size
- Collect more edge cases
- Fine-tune deeper layers of InceptionV3
- Implement confidence threshold for unknown detection
- Explore other pre-trained models (EfficientNet, MobileNet)

Conclusion & Key Findings



- Developed an Instant Cup Noodles Condition Detector (Finished, Opened, Sealed)
- Transfer learning (InceptionV3) improved accuracy and training efficiency
- Larger image sizes (224×224 / 299×299) improved feature extraction
- Overfitting was reduced using data augmentation, dropout, batch normalization, and early stopping
- Video-to-image extraction increased dataset diversity and robustness
- Final model shows high test accuracy and good generalisation, making it suitable for real-world use

THANK YOU